

Mannan Shukla

Sunnyvale, California • US Citizen

(609) 575-0436 • mannanshukla@icloud.com • [linkedin.com/in/mannanshukla](https://www.linkedin.com/in/mannanshukla) • github.com/mannanshukla

TECHNICAL SKILLS

Languages: Python, Go, Bash, TypeScript

Inference: vLLM, SGLang, TensorRT, NVIDIA Dynamo, CUDA, NCCL, PagedAttention, KV cache, continuous batching

ML & Agents: PyTorch, Hugging Face, LoRA, LangGraph, MCP, Kubeflow

Infrastructure: Kubernetes, Docker, Helm, Terraform, AWS, ArgoCD, Cilium, Slurm

Production: Postgres (Neon), Caddy, Stripe, Auth0, Prometheus, Grafana, GitHub Actions, CI/CD, GitOps

EXPERIENCE

Founder & Engineer

spawnpoint • San Francisco, CA

July 2026 – Present

- Building **spawnpoint** (getspawnpoint.com): an AI agent deploys an app over **MCP** and gets a public HTTPS link
- Built the **MCP** server in **Go** (23 tools, stateless HTTP) and an **OAuth 2.1** server: **PKCE**, one-shot codes, rotation
- Replaced SSH pushes with a **pull agent** on each VM (cloud-init bootstrap, one-time tokens, long-poll jobs, presigned **S3** bundles): 4 MB deploys 93 s to 72 s; four runtimes incl. **Docker**
- Shipped gated sharing (email-OTP viewers, allow lists), DNS-verified **custom domains**, scale-to-zero on idle
- Built per-project managed **Postgres** and **S3** storage, **Stripe** billing, passwordless **Auth0**, sealed env vars
- Run production on **AWS** (EC2, **RDS**, **S3**) with zero-downtime deploys over **SSM** from **GitHub Actions**

Software Engineer

Voltage Park / Lightning AI • San Francisco, CA & Redmond, WA

July 2025 – July 2026

- Shipped hot model-weight reloading for **vLLM** inference deployments via a **Go** sidecar, cutting fine-tuned model rollout time from 10–30 minutes to seconds without restarting GPU workers
- Built core pieces of an **S3**-compatible object store (**VAST**-backed): a health-aware **P2C** client-side load balancer over CNode IPs with lock-free in-flight tracking and automatic failover, plus **Terraform/Vault**-managed **RDS** infra
- Built distributed inference orchestration in **Go** on **Kubernetes**, implementing deployment-target controllers and **DynamoDB** event sourcing to reconcile and recover fleets of GPU-serving replicas
- Created the **vp-cloud-sdk Python** SDK for Voltage Park's orchestration platform, exposing fleet/node/remediation APIs and a full staging release pipeline with **GitHub Actions**
- Own a multi-tenant GPU platform spanning **ArgoCD**, **Terraform**, **Helm**, **Cilium/Envoy**, NVIDIA GPU Operator, and **Prometheus/Grafana**, with tenant isolation and production observability
- Built a **Slurm** GPU health-check suite for distributed training (ECC, PCIe AER, **NCCL** stalls, **DCGM** checks, thermal) that auto-drains unhealthy nodes, creates RMA tickets, and sends **Slack** alerts to protect multi-node jobs
- Operate production **H100** clusters, resolving incidents across GPU scheduling, distributed networking, and **NCCL** workloads

Undergraduate Researcher: AIOS

Rutgers University • New Brunswick, NJ

Apr 2024 – Apr 2025

- Built an **LLM** agent recommender hitting 83% accuracy vs. **GPT-4** with a 1B model via **LoRA** fine-tuning (**Unsloth**)

Partnerships Intern

Nas.io • New York, NY

Feb 2025 – May 2025

- Built an **Instagram scraper** and **Python/Klaviyo** pipeline that lifted outreach 50x; led partnerships to \$2K GMV

PROJECTS

BYOA: Bring Your Own Agent

- Built a **Kubernetes** inference platform that deploys agents on arbitrary **Hugging Face** models behind OpenAI-compatible APIs, using **NVIDIA Dynamo** for disaggregated serving, KV-cache-aware routing, and GPU autoscaling

TaskCLI: SWE-bench-style evaluation runner

- Built a **Python** agent-evaluation harness for SWE-bench-style task bundles with reproducible **Docker** sandboxes, hidden-test patch grading, **JUnit XML** parsing, **SQLite** artifact tracking, and OpenAI-compatible solvers

tochi.one: AI property management agents

- Built payment, lease-generation, and tenant-communication workflows in **LangGraph/MCP** with **Stripe** and **Twilio**

AdBlockIRL: AR ad-blocking

- Trained **YOLOv8** on **Roboflow** to detect 100+ concurrent ads at 5ms latency, applying real-time mosaic censoring via **OpenCV**; won \$2000 at Bain Capital Ventures AI Hackathon

EDUCATION

Rutgers University – New Brunswick

Bachelor of Science – Computer Science • Cum Laude • Linux User Group (President)

May 2025