

Mannan Shukla

Sunnyvale, California • US Citizen

(609) 575-0436 • mannanshukla@icloud.com • [linkedin.com/in/mannanshukla](https://www.linkedin.com/in/mannanshukla) • github.com/mannanshukla

TECHNICAL SKILLS

Languages: Python, Go, Bash, TypeScript

Inference: vLLM, SGLang, TensorRT, NVIDIA Dynamo, CUDA, NCCL, PagedAttention, KV cache, continuous batching

ML & Agents: PyTorch, Hugging Face, LoRA, LangGraph, MCP, Kubeflow

Infrastructure: Kubernetes, Docker, Helm, Terraform, AWS, ArgoCD, Cilium, Slurm

Production: Postgres (Neon), Caddy, Stripe, Auth0, Prometheus, Grafana, GitHub Actions, CI/CD, GitOps

EXPERIENCE

Founder & Engineer

spawnpoint • San Francisco, CA

July 2026 – Present

- Building **spawnpoint** (getspawnpoint.com), an agent-native hosting platform: an AI agent deploys an app over **MCP** and gets back a public HTTPS link, private by default and shared like a doc
- Built the **MCP** server in **Go** (17 tools, stateless Streamable HTTP) and the **OAuth 2.1** server behind it: dynamic client registration, **PKCE** S256, single-use codes, rotating refresh tokens with reuse detection
- Built the deploy path: boot a CPU VM, push the bundle over **SSH** in chunked commands with an atomic swap, retry through a lossy proxy, reconcile orphans at startup; four runtimes (static, Node, Python, **Docker**), builds run on the user's own VM
- Shipped gated sharing: a host-routed share gate with email-OTP viewers and per-request allow lists, DNS-verified **custom domains** with on-demand **TLS** via **Caddy**, and scale-to-zero that stops idle VMs and wakes them on the next visit
- Built **Stripe** subscription billing (hosted checkout, portal, signature-verified webhooks), passwordless **Auth0** sign-in, stateless **HMAC**-signed sessions, and sealed-at-rest per-project env vars
- Run production with **Postgres** on **Neon**, zero-downtime blue/green deploys behind Caddy, and **GitHub Actions** CI/CD (gofmt, vet, race tests, doc-drift checks); ship a **Claude Code** plugin auto-published to a public marketplace

Software Engineer

Voltage Park / Lightning AI • San Francisco, CA & Redmond, WA

July 2025 – July 2026

- Shipped hot model-weight reloading for **vLLM** inference deployments via a **Go** sidecar, cutting fine-tuned model rollout time from 10–30 minutes to seconds without restarting GPU workers
- Built core pieces of an **S3**-compatible object store (**VAST**-backed): a health-aware **P2C** client-side load balancer over CNode IPs with lock-free in-flight tracking and automatic failover, plus **Terraform/Vault**-managed **RDS** infra
- Built distributed inference orchestration in **Go** on **Kubernetes**, implementing deployment-target controllers and **DynamoDB** event sourcing to reconcile and recover fleets of GPU-serving replicas
- Created the **vp-cloud-sdk Python** SDK for Voltage Park's orchestration platform, exposing fleet/node/remediation APIs and a full staging release pipeline with **GitHub Actions**
- Own a multi-tenant GPU platform spanning **ArgoCD**, **Terraform**, **Helm**, **Cilium/Envoy**, NVIDIA GPU Operator, and **Prometheus/Grafana**, with tenant isolation and production observability
- Built a **Slurm** GPU health-check suite for distributed training (ECC, PCIe AER, **NCCL** stalls, **DCGM** checks, thermal) that auto-drains unhealthy nodes, creates RMA tickets, and sends **Slack** alerts to protect multi-node jobs
- Operate production **H100** clusters, resolving incidents across GPU scheduling, distributed networking, and **NCCL** workloads

Undergraduate Researcher: AIOS

Rutgers University • New Brunswick, NJ

Apr 2024 – Apr 2025

- Built an **LLM** agent recommender hitting 83% accuracy vs. **GPT-4** with a 1B model via **LoRA** fine-tuning (**Unsloth**)

Partnerships Intern

Nas.io • New York, NY

Feb 2025 – May 2025

- Built an **Instagram scraper** and **Python/Klaviyo** pipeline that lifted outreach 50x; led partnerships to \$2K GMV

PROJECTS

BYOA: Bring Your Own Agent

- Built a **Kubernetes** inference platform that deploys agents on arbitrary **Hugging Face** models behind OpenAI-compatible APIs, using **NVIDIA Dynamo** for disaggregated serving, KV-cache-aware routing, and GPU autoscaling

TaskCLI: SWE-bench-style evaluation runner

- Built a **Python** agent-evaluation harness for SWE-bench-style task bundles with reproducible **Docker** sandboxes, hidden-test patch grading, **JUnit XML** parsing, **SQLite** artifact tracking, and OpenAI-compatible solvers

tochi.one: AI property management agents

- Built payment, lease-generation, and tenant-communication workflows in **LangGraph/MCP** with **Stripe** and **Twilio**

AdBlockIRL: AR ad-blocking

- Trained **YOLOv8** on **Roboflow** to detect 100+ concurrent ads at 5ms latency, applying real-time mosaic censoring via **OpenCV**; won \$2000 at Bain Capital Ventures AI Hackathon

EDUCATION

Rutgers University – New Brunswick

Bachelor of Science – Computer Science • Cum Laude • Linux User Group (President)

May 2025